

Cornell University
Cornell Tech,
New York, NY

arush@cornell.edu
<http://rush-nlp.com>
@srush_nlp

Alexander M. Rush

Appointment	<i>Cornell University Computer Science Department</i> 2019-. Associate Professor of Computer Science (without Tenure)	
	<i>Hugging Face</i> 2019-. NLP Open-Source Startup	
	<i>Harvard University School of Engineering and Applied Sciences</i> 2015-2019. Assistant Professor of Computer Science	
	<i>Facebook Artificial Intelligence Research Lab</i> 2015. Post-Doctoral Fellowship Advisor: Yann LeCunn	
Education	<i>Massachusetts Institute of Technology</i> 2009-2014. Ph.D, Computer Science. Advisor: Michael Collins Dissertation: <i>Relaxation Methods for Natural Language Decoding.</i>	
	<i>Harvard University</i> 2007. B.A., Computer Science.	
Grants, Awards, and Major Service	2020	Best Demo Paper - Runner-Up, ACL Best Paper - DAC (Hardware)
	2019	NSF Career Award Sony Faculty Awards General Chair and Senior Program Chair , ICLR Best Demo Paper - Nominee, ACL
	2018	Senior Program Chair, ICLR Google, Facebook, and Amazon AWS Faculty Awards Best Paper - Runner-Up, VAST (Visualization) Keynote, AMTA (American Machine Translation Association)
	2017	Bloomberg and Intel AI Collaboration Faculty Awards Best Demo - Runner-Up, ACL
	2016	Microsoft Azure and Samsung AI Award Invitation IJCAI Early Research Spotlight Best Paper - Runner-Up, EMNLP
	2015	NIPS Deep Learning Symposium (Invited Paper) Google Faculty Award
	2012	Best Paper Award, NAACL
	2010	Best Paper Award, EMNLP

Publications

(Full list: <http://bit.do/alexander-rush>)

Highly Cited Publications (Google Scholar Metrics)

Alexander M. Rush, Sumit Chopra, and Jason Weston. *A Neural Attention Model for Abstractive Sentence Summarization*. EMNLP 2015. (1600 citations, 3rd Most Cited AAAI Paper 2015-2020)

Yoon Kim, Yacine Jernite, David Sontag, and Alexander M. Rush. *Character-Aware Neural Language Models*. AAAI 2016, (1250 citations, 3rd Most Cited AAAI Paper 2015-2020)

Hendrik Strobelt, Sebastian Gehrmann, Hanspeter Pfister, and Alexander M. Rush. *LSTMVis: A Tool for Visual Analysis of Hidden State Dynamics in Recurrent Neural Networks*. InfoVis 2017, (139 citations, 12th Most Cited IEEE Transactions on Visualization and Computer Graphics Paper 2015-2020)

Guillaume Klein, Yoon Kim, Yuntian Deng, Jean Senellart, Alexander M. Rush. *OpenNMT: Open-Source Toolkit for Neural Machine Translation*. ACL Demo 2017 (900 citations, Best Demo Runner-up, 5th Most Cited ACL Paper 2015-2020)

All Conference Papers

- [1] Thierry Tambe, En-Yu Yang, Zishen Wan, Yuntian Deng, Vijay Janapa Reddi, Alexander Rush, David Brooks, Gu-Yeon Wei. *AdaptiveFloat: A Floating-point based Data Type for Resilient Deep Learning Inference*. DAC 2020 (Best Paper)
- [2] Thomas Wolf et al. *Transformers: State-of-the-art Natural Language Processing*. EMNLP Demos 2020
- [3] Yuntian Deng, Alexander M. Rush. *Cascaded Text Generation with Markov Transformers*. Preprint
- [4] Victor Sanh, Thomas Wolf, Alexander M. Rush. *Movement Pruning: Adaptive Sparsity by Fine-Tuning*. Preprint
- [5] Alexander Rush. *Torch-Struct: Deep Structured Prediction Library*. ACL Demos 2020 (Best Demo Honorable Mention)
- [6] Noriyuki Kojima, Hadar Averbuch-Elor, Alexander M. Rush, Yoav Artzi. *What is Learned in Visually Grounded Neural Syntax Acquisition*. ACL 2020 (Short)
- [7] Xiang Lisa Li, Alexander M. Rush. *Posterior Control of Blackbox Generation*. ACL 2020
- [8] Jiawei Zhou, Zhiying Xu, Alexander M. Rush, Minlan Yu. *Automating Botnet Detection with Graph Neural Networks*. AutoML for Networking and Systems Workshop

- [9] Georgios A. Tritsarlis, Yiqi Xie, Alexander M. Rush, Stephen Carr, Marios Mattheakis, Efthimios Kaxiras. *LAN – A materials notation for 2D layered assemblies*. None
- [10] Udit Gupta, Brandon Reagen, Lillian Pentecost, Marco Donato, Thierry Tambe, Alexander M. Rush, Gu-Yeon Wei, David Brooks. *MASR: A Modular Accelerator for Sparse RNNs*. PACT 2019
- [11] Joe Davison, Joshua Feldman and Alexander Rush. *Commonsense Knowledge Mining from Pretrained Models*. EMNLP 2019
- [12] Zachary Ziegler, Yuntian Deng and Alexander Rush. *Neural Linguistic Steganography*. EMNLP 2019
- [13] Yoon Kim, Chris Dyer, Alexander M. Rush. *Compound Probabilistic Context-Free Grammars for Grammar Induction*. ACL 2019
- [14] Gehrmann S, Strobel H, Krueger R, Pfister H, and Alexander M. Rush. *Visual Interaction with Deep Learning Models through Collaborative Semantic Inference*. InfoVis 2019
- [15] Jiawei Zhou, Alexander M. Rush. *Simple Unsupervised Summarization by Contextual Matching*. ACL 2019
- [16] Sebastian Gehrmann, Hendrik Strobel, Alexander M. Rush. *GLTR: Statistical Detection and Visualization of Generated Text*. ACL Demo 2019 (Best Demo Honorable Mention)
- [17] Yoon Kim, Alexander M. Rush, Lei Yu, Adhiguna Kuncoro, Chris Dyer, Gabor Melis. *Unsupervised Recurrent Neural Network Grammars*. NAACL 2019
- [18] Adji B. Dieng, Yoon Kim, Alexander M. Rush, David M. Blei. *Avoiding Latent Variable Collapse With Generative Skip Models*. AISTATS 2019
- [19] Fritz Obermeyer, Eli Bingham, Martin Jankowiak, Justin Chiu, Neeraj Pradhan, Alexander Rush, Noah Goodman. *Tensor Variable Elimination for Plated Factor Graphs*. ICML 2019
- [20] Zachary M. Ziegler, Alexander M. Rush. *Latent Normalizing Flows for Discrete Sequences*. ICML 2019
- [21] Yoon Kim, Sam Wiseman, Alexander M. Rush. *Deep Latent-Variable Models for Natural Language*. EMNLP 2018 (Tutorial)
- [22] Sebastian Gehrmann, Falcon Z. Dai, Henry Elder, Alexander M. Rush. *End-to-End Content and Plan Selection for Data-to-Text Generation*. INLG 2018
- [23] Yuntian Deng*, Yoon Kim*, Justin Chiu, Demi Guo, Alexander M. Rush. *Latent Alignment and Variational Attention*. NIPS 2018
- [24] Sam Wiseman, Stuart M. Shieber, Alexander Rush. *Learning Neural Templates for Text Generation*. EMNLP 2018

- [25] Sebastian Gehrmann, Yuntian Deng, Alexander Rush. *Bottom-Up Abstractive Summarization*. EMNLP 2018
- [26] Luke Melas-Kyriazi, George Han, Alexander Rush. *Training for Diversity in Image Paragraph Captioning*. EMNLP 2018 (Short)
- [27] Luong Hoang, Sam Wiseman, Alexander Rush. *Entity Tracking Improves Cloze-style Reading Comprehension*. EMNLP 2018 (Short)
- [28] Hendrik Strobelt, Sebastian Gehrmann, Michael Behrisch, Adam Perer, Hanspeter Pfister, Alexander M. Rush. *Seq2Seq-Vis: A Visual Debugging Tool for Sequence-to-Sequence Models*. VAST 2018, EMNLP-BlackBox 2018 (Best Paper - Honorable Mention)
- [29] Alexander M. Rush. *The Annotated Transformer*. ACL NLP-OSS 2018
- [30] Jean Senellart, Dakun Zhang, Bo Wang, Guillaume Klein, J.P. Ramachandirin, Josep Crego, Alexander M. Rush. *OpenNMT System Description for WNMt 2018: 800 words/sec on a single-core CPU*. WNMt 2018 (First-Place CPU Speed/Memory)
- [31] Yoon Kim, Sam Wiseman, Andrew C. Miller, David Sontag, Alexander M. Rush. *Semi-Amortized Variational Autoencoders*. ICML 2018
- [32] Brandon Reagen, Udit Gupta, Robert Adolf, Michael M. Mitzenmacher, Alexander M. Rush, Gu-Yeon Wei, David Brooks. *Compressing Deep Neural Networks with Probabilistic Data Structures*. ICML 2018, SysML 2018
- [33] Allen Schmaltz, Yoon Kim, Alexander M. Rush, Stuart M. Shieber. *Adapting Sequence Models for Sentence Correction*. EMNLP 2017
- [34] Sam Wiseman, Stuart M Shieber Alexander M. Rush. *Challenges in Data-to-Document Generation*. EMNLP 2017
- [35] Junbo Zhao, Yoon Kim, Kelly Zhang, Alexander M. Rush, Yann LeCun. *Adversarially Regularized Autoencoders*. ICML 2018, NIPS 2017 Workshop
- [36] Guillaume Klein, Yoon Kim, Yuntian Deng, Jean Senellart, Alexander M. Rush. *OpenNMT: Open-Source Toolkit for Neural Machine Translation*. ACL Demo 2017 (Best Demo Runner-up)
- [37] Ankit Gupta, Alexander M. Rush. *Dilated Convolutions for Modeling Long-Distance Genomic Dependencies*. ICML CompBio 2017 (Best Poster)
- [38] Yuntian Deng, Anssi Kanervisto, Jeffrey Ling, and Alexander M. Rush. *Image-to-Markup Generation with Coarse-to-Fine Attention*. ICML 2017
- [39] Hendrik Strobelt, Sebastian Gehrmann, Hanspeter Pfister, and Alexander M. Rush. *LSTMVis: A Tool for Visual Analysis of Hidden State Dynamics in Recurrent Neural Networks*. InfoVis 2017
- [40] Yoon Kim, Carl Denton, Luong Hoang, and Alexander M. Rush. *Structured Attention Networks*. ICLR 2017

- [41] Greg Yang and Alexander M. Rush. *Lie-Access Neural Turing Machines*. ICLR 2017
- [42] Yoon Kim and Alexander M. Rush. *Sequence-Level Knowledge Distillation*. EMNLP 2016
- [43] Sam Wiseman and Alexander M. Rush. *Sequence-to-Sequence Learning as Beam-Search Optimization*. EMNLP 2016 (Best Paper Runner-Up)
- [44] Peter Kraft, Hirsh Jain, and Alexander M. Rush. *An Embedding Model for Predicting Roll-Call Votes*. Proceedings of EMNLP 2016
- [45] Allen Schmaltz, Alexander M. Rush, and Stuart M. Shieber. *Word Ordering Without Syntax*. EMNLP 2016
- [46] Allen Schmaltz, Yoon Kim, Alexander M. Rush, and Stuart M. Shieber. *Sentence-Level Grammatical Error Identification as Sequence-to-Sequence Correction*. Workshop Submission for AESW 2016 (Top Performing System)
- [47] Sam Wiseman, Alexander M. Rush, and Stuart M. Shieber. *Learning Global Features for Coreference Resolution*. NAACL 2016
- [48] Sumit Chopra, Michael Auli, and Alexander M. Rush. *Abstractive Sentence Summarization with Attentive Recurrent Neural Networks*. NAACL 2016
- [49] Yoon Kim, Yacine Jernite, David Sontag, and Alexander M. Rush. *Character-Aware Neural Language Models*. AAAI 2016
- [50] Alexander M. Rush, Sumit Chopra, and Jason Weston. *A Neural Attention Model for Abstractive Sentence Summarization*. EMNLP 2015.
- [51] Jason Weston, Antoine Bordes, Sumit Chopra, Tomas Mikolov, and Alexander M. Rush. *Towards AI-Complete Question Answering A Set of Prerequisite Toy Tasks*. ArXiv Preprint
- [52] Sam Wiseman, Alexander M. Rush, Jason Weston, and Stuart M. Shieber. *Learning Anaphoricity and Antecedent Ranking Features for Coreference Resolution*. ACL 2015.
- [53] Yacine Jernite, Alexander M. Rush, and David Sontag. *A Fast Variational Approach for Learning Markov Random Field Language Models*. ICML 2015.
- [54] Lingpeng Kong, Alexander M. Rush, and Noah A. Smith. *Transforming Dependencies into Phrase Structures*. NAACL 2015.
- [55] Yin-Wen Chang, Alexander M. Rush, John DeNero, and Michael Collins.. *A Lagrangian Relaxation Algorithm for Bilingual Word Alignment*. Proceedings of ACL 2014.
- [56] Alexander M. Rush, Yin-Wen Chang, and Michael Collins.. *Optimal Beam Search for Machine Translation*. Proceedings of EMNLP 2013.
- [57] Karl Stratos, Alexander M. Rush, Shay B. Cohen, and Michael Collins.. *Spectral Learning of Refinement HMMs*. Proceedings of CoNLL 2013.

- [58] Alexander M. Rush, Roi Reichert, Michael Collins, and Amir Globerson.. *Improved Parsing and POS Tagging Using Inter-Sentence Consistency Constraints*. *Proceedings of EMNLP 2012*.
- [59] Alexander M. Rush and Slav Petrov. *Vine Pruning for Efficient Multi-Pass Dependency Parsing*. *Proceedings of NAACL 2012*. (Best Paper Award)
- [60] Alexander M. Rush and Michael Collins.. *Exact Decoding of Syntactic Translation Models through Lagrangian Relaxation*. *Proceedings of ACL 2011*.
- [61] Terry Koo, Alexander M. Rush, Michael Collins, Tommi Jaakkola, and David Sontag.. *Dual Decomposition for Parsing with Non-Projective Head Automata*. *Proceedings of EMNLP 2010*.
- [62] Alexander M. Rush, David Sontag, Michael Collins, and Tommi Jaakkola. *On Dual Decomposition and Linear Programming Relaxations for Natural Language Processing*. *Proceedings of EMNLP 2010*.
- [63] Rebecca Nesson, Stuart M. Shieber, and Alexander M. Rush.. *Induction of probabilistic synchronous tree-insertion grammars for machine translation*. *Proceedings of AMTA 2006*.

Journal Papers

Alexander M. Rush and Michael Collins. *A Tutorial on Dual Decomposition and Lagrangian Relaxation for Inference in Natural Language Processing*. Journal of Artificial Intelligence Research 2012.

Service

General Chair: ICLR (Deep Learning) 2020
 Senior Program Chair: ICLR (Deep Learning) 2019
 Senior Area Chair: NeurIPS 2019; ACL (Generation) 2019; ACL (Generation) 2020
 Area Chair: NAACL (Machine Learning) 2020; NIPS 2018; NAACL (Machine Learning) 2016; ICLR (NLP) 2017, 2018; EMNLP (Generation/Summarization) 2017; ACL (Parsing and Tagging) 2017
 Tutorial Chair: NAACL 2016

Teaching

- 2020 Instructor, Topic in Machine Learning and NLP , Cornell Tech, Spring.
- 2019 Instructor, Machine Learning for NLP , Harvard University, Spring.
- 2018 Instructor, Machine Learning for NLP (Rated 4.8/5) , Harvard University, Spring.
Instructor, Advanced Machine Learning , Harvard University, Fall.
- 2017 Instructor, Machine Learning , Harvard University, Spring.
Instructor, Graduate Machine Learning , Harvard University, Fall.
- 2016 Instructor, Machine Learning for NLP (Rated 4.9/5), Harvard University, Spring.
- 2015 Instructor, Artificial Intelligence, Harvard University, Fall.
- 2013 Instructor (with Michael Collins), Natural Language Processing, Columbia University, Fall.
Head Teaching Assistant, Natural Language Processing , Michael Collins, Columbia University, Spring (taught on Coursera, 30,000+ registered students).
- 2012 Head Teaching Assistant, Natural Language Processing, Michael Collins, Columbia University, Fall.

Patents

- A neural attention model for abstractive summarization (Facebook). Alexander M. Rush, Sumit Chopra, Jason Weston 2017.
- Techniques for discriminative dependency parsing (Google). Slav Petrov, Alexander M. Rush, 2015.
- Efficient parsing with structured prediction cascades (Google). Slav Petrov, Alexander M. Rush, 2013
- Determining user affinity towards applications on a social networking website (Facebook., Thomas S. Whitnah, Alexander M. Rush, Ding Zhou, Ruchi Sangvhi, 2010.

Personal Libraries

[MiniConf.](#)

MiniConf is a project developed for ICLR as an easy-to-use tool for hosting fully remote asynchronous virtual conferences. It was heavily used in 2020 to host ACL, ICML, AKBC, AISTATS, and many other virtual conferences..

[NLP Browser.](#)

NLP Browser is a web app that lets any easily browse through more than 150 datasets used in NLP and hosted by Hugging Face. The app is a pretty addictive way to casually learn about new datasets and challenges..

NamedTensor (Tensor Considered Harmful).

Named Tensor is a proposal for adding a new datastructure to mathematical libraries to treat tensors more like dicts and less like tuples. This blog post had the impact of getting PyTorch to add a NamedTensor annotation in v1.3 of the library..

Torch Struct.

Torch-Struct is a passion project of mine to test out whether deep learning libraries can be used to implement classical structured prediction. It includes heavily-tested reference reimplementations of many core NLP algorithms..

The Annotated Transformer.

The annotated transformer was an experiment in blogging based on literature papers. The idea was to teach researchers how an important model in NLP works by aligning the paper line-by-line with an implementation. The blog post was widely distributed, and there have been many follow-ups for new model..

Academic Internships

Research Intern, *Google Research*, 2011 – 2013 , New York, NY. Advisor: Slav Petrov.
Research Intern, *USC/ISI* , Summer 2010, Marina Del Rey, CA. Advisor: Liang Huang.

Industry

Lead Engineer (Platform Team), *Facebook*, 2007 – 2009, Palo Alto, CA.
Developed compiler for Facebook Markup Language (FBML) to sanitize user content.
Developed system for crowd-sourced translation of Facebook user text.

Invited Talks

2019 Colloquium, University of Edinburgh, Spring.
Colloquium, Tel Aviv University, Spring.
Invited Talk, Conversational Intelligence Summer.
Invited Talk, GANocracy, MIT, Summer.
Invited Talk, OpenAI, MIT, Summer.
Invited Talk, NeuralGen Workshop NAACL Summer.
Invited Talk, Berkeley NLP, Fall.
Keynote, PyTorch Developers Conference, Fall.

- 2018 Invited Talk, University of Washington, Spring.
 Invited Talk, Allen Institute for AI, Spring.
 Invited Talk, MSR, Spring.
 Keynote, American Machine Translation Association, Spring.
 Invited Talk, University of Texas, Spring.
 Invited Talk, University of Maryland, Spring.
 Invited Talk, Georgetown, Spring.
 Invited Talk, Lisbon ML Summer School, Summer.
 Invited Talk, Columbia University, Fall.
 Invited Talk, New York University - Text as Data, Fall.
 Tutorial, EMNLP, Fall.
- 2017 Invited Talk, Google Faculty Day, Spring.
 Invited Talk, New England Machine Learning Day, Spring.
 Invited Talk, Google, Spring.
 Invited Talk, Berkeley CS, Spring.
 Invited Talk, Notre Dame, Spring.
 Colloquium, TTI-Chicago, Spring.
 Invited Talk, Apple, Siri Team, Spring.
 Colloquium, Samsung Global AI Forum, Fall.
 Invited Talk, AMD, Fall.
- 2016 Invited Talk, NYU, Fall.
 Invited Talk, BBN Research, Fall.
 Invited Talk, Bloomberg, Fall.
 Invited Talk (Speech Group), MIT, Fall.
 Invited Talk, IBM Research, Fall.
 Colloquium, CMU, Fall.
 Invited Talk, Stanford NLP, Summer.
 Invited Talk, Oracle Labs, Summer.
 Invited Talk, Twitter, Summer.
 Colloquium, John Hopkins University, Spring.
 Colloquium, Rakuten, Spring.

- 2014 Colloquium, University of Washington, Spring.
 Colloquium, NYU, Spring.
 Colloquium, CMU, Spring.
 Colloquium, MIT, Spring.
 Colloquium, Harvard, Spring.
 Colloquium, TTIC, Spring.
 Colloquium, University of Maryland, Spring.
- 2013 Invited Tutorial, UMBC, October.
 Invited Talk, CS and Social Science Seminar, UMass Amherst, October.
 Talk, NLP Seminar, Columbia University, October.
 Invited Talk, ML Seminar, UMass Amherst, October.
 Invited Talk, Johnson Research Labs, NY, August.
 Invited Talk, Society for Historians of American Foreign Relations, Arlington, June.
 Invited Talk, Columbia University, Spring.
 Invited Talk, NLP Seminar, City University of New York, Spring.
- 2012 Invited Tutorial, Neural Information Processing Systems (NIPS), December.
- 2011 Invited Tutorial, Google Research, Mountain View, August.
 Tutorial. Association of Computational Linguistic (ACL), June.
 Invited Talk, ML Seminar, University of Massachusetts, Amherst, Spring.
 ML Tea, MIT, January.
- 2010 NLP Seminar, USC/ISI, Summer.
- 2006 Invited Talk, Computational Linguistics Seminar, University of Pennsylvania, November.