Cascaded Text Generation with Markov Transformers

Yuntian DENG, Alexander M. RUSH
Harvard University, Cornell University

github.com/harvardnlp/cascaded-generation

We present a **fast and fluent** text generation method using **bounded-order** Markov models. To decode from high-order models efficiently, we propose a cascaded decoding approach that **prunes the search space** using lower-order models, and we introduce a Markov transformer that can **parameterize the entire cascade**.

1/ Motivation

Fully Autoregressive (AR)

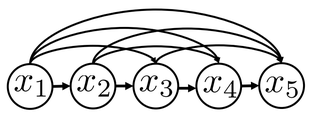

- Decoding: beam search
- Fluent but serial

Nonautoregressive (NAR)

- Decoding: argmax
- Parallel but disfluent

- We propose a method that's both fast and fluent

2/ Model: Bounded-Order MRF

- An m-th order Markov Random Field (MRF)

$$P^{(m)}(x_{1:L};\theta) \propto \exp \sum_{l=1}^{L-m} f_l^{(m)}(x_{l:l+m};\theta)$$

- Each $f_l^{(m)}$ models dependencies among adjacent $(m+1)$ words
- Example: $m=1$, $L=5$

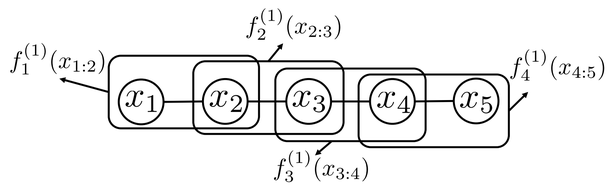

- Special Cases:
  - $m=0$: nonautoregressive
  - $m=\mathrm{L}-1$: fully autoregressive
  - $0<m<\mathrm{L}-1$: **bounded-order models** (this work)

3/ Parallel Decoding

- Bounded-order models can be decoded in parallel
- In practice, takes <1% total time

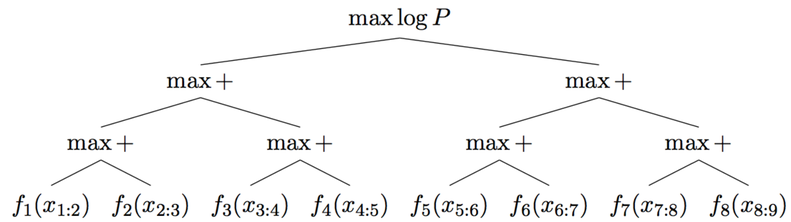

4/ Cascaded Decoding

- Parallel decoding for an M-th order model takes $O(V^{M+1}\log L)$ time
- Impractical even when M=1
- Cascaded decoding filters unlikely n-grams using lower-order models
- Example: M= 3

a) filter unlikely unigrams using $m=0$

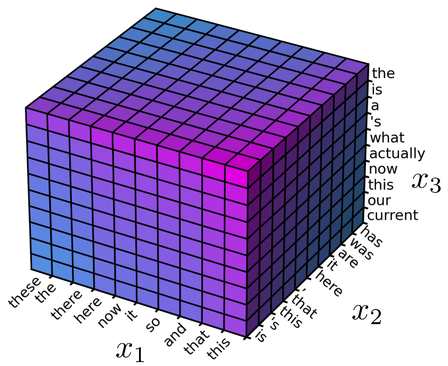

b) filter unlikely bigrams using $m=1$

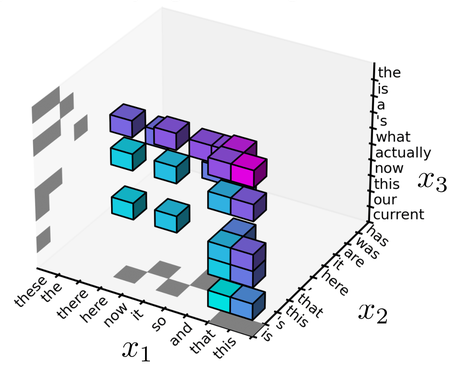

c) filter unlikely trigrams using $m=2$

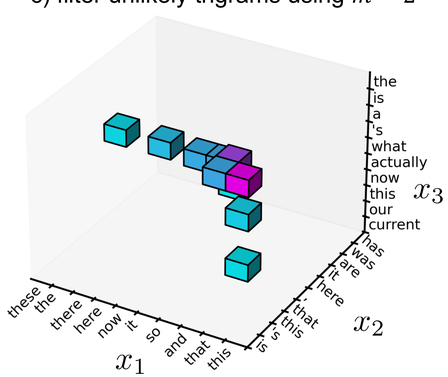

d) parallel decode in the pruned search space

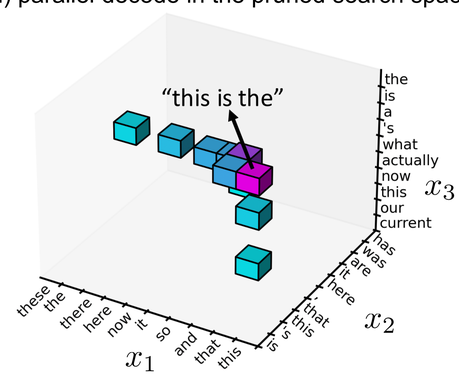

5/ Parameterization: Markov Transformer

- Parameterize the entire cascade with a single transformer
- Training: insert M-spaced barriers with random offset

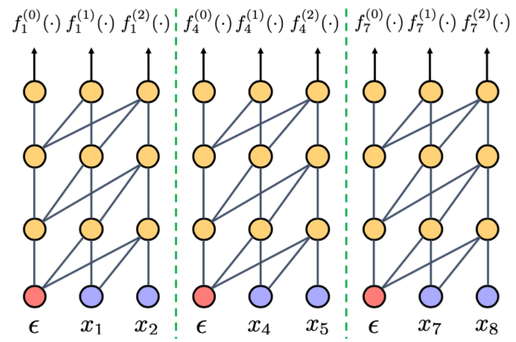

- Test: can be applied as any Markov model with $m < M$

6/ Results & Analysis

- Speed/accuracy tradeoff

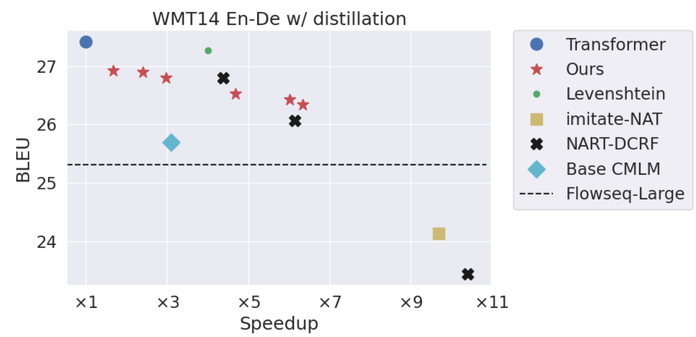

- Fewer repetitions

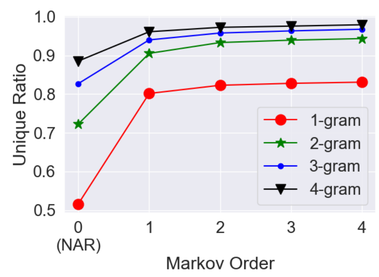

- More hypotheses scored